
A Probing Confound That Survives the Untrained-Network Control

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The standard defence of a probing result is a floor: if an untrained network supports
2 the same probe far less well, the probe is reading something the model learned. We
3 report a case where that defence passes and the conclusion is still mostly wrong. A
4 linear probe on a code model’s residual stream appears to recover variable roles
5 language-independently, moving -0.012 macro-F1 across a boundary where a
6 model-free surface classifier loses 0.126, and it clears an untrained network by
7 $+0.269$. But the probe pools over the occurrence’s own tokens, so its input contains
8 the variable’s name, and names are among the most language-invariant features
9 of parallel corpora. Excluding the name from the pooling, with the forward pass
10 unchanged, removes 61% of the probe’s margin over a matched untrained floor
11 and multiplies its boundary effect fivefold, leaving it no flatter than the untrained
12 network itself. The untrained control cannot catch this confound, because the
13 untrained network reads the same name and inherits the same invariance. What
14 caught it was a model-free comparator that could not see the name. We suggest
15 such comparators belong beside untrained floors in the standard battery, and report
16 a negative result for role probing: with the identifier excluded, we find no evidence
17 of language-independent role representation at this scale.

18 1 Introduction

19 Probing classifiers come with a known epistemic gap: a probe reports what is decodable, not what
20 the model uses, and a sufficiently expressive probe can decode from noise [Belinkov, 2022, Hewitt
21 and Liang, 2019]. The field’s working answer is a battery of controls, and the strongest of them is the
22 untrained network: run the same probe on the same architecture with random weights, and whatever
23 survives is at least *learned*, even if not used [Zhang and Bowman, 2018, Voita and Titov, 2020].

24 This paper is about a failure mode that battery does not cover. The controls above all vary the *model*
25 while holding the probe’s input fixed. If the input itself carries the label, every member of the battery
26 inherits the signal: the trained network decodes it, and the untrained network decodes it too, so the
27 floor rises with the effect and the comparison stays flattering. Nothing in the standard battery varies
28 what the probe is allowed to read.

29 We hit this in a concrete setting: cross-lingual transfer of variable-role probes in code models. The
30 confounded conclusion was attractive, plausible, and supported by an untrained floor. It fell to a
31 control from outside the standard battery, and the case is clean enough to serve as a template for
32 where else the same gap may hide.

Table 1: Cross-lingual role transfer, macro-F1 over three roles. *Close* is JavaScript $\leftrightarrow$ PHP; *Python* is the four transfers involving Python. Each probe is paired with an untrained network at its own pooling, because different poolings read different token sets and have different achievable floors.

	close	Python	effect
Surface n -gram, name masked (no model)	0.929	0.804	-0.126
Probe, span-pooled (reads the name)	0.852	0.840	-0.012
untrained, span-pooled [†]	0.590	0.568	-0.022
Probe, context-pooled (no name)	0.630	0.569	-0.061
untrained, context-pooled	0.527	0.464	-0.064

[†] Committed, but measured on an earlier, larger sample rather than the three-way intersection the other rows use.




Figure 1: Table 1 as a picture. The surface classifier tracks the language boundary. The span-pooled probe does not, but its input contains the variable’s name. Excluding the name drops the probe to just above its untrained floor, and both slope downward together.

2 The setting, and the seductive result

A *variable role* [Sajaniemi, 2002], such as accumulator or index key, describes what a variable does rather than what it is called, so it is a natural candidate for language-independent representation. XLCoST [Zhu et al., 2022] provides the same problems solved in Python, JavaScript and PHP; we restrict to the 869 problems all three share, so every comparison sees identical problem ids, and fit a role classifier on one language while testing on another, over all six ordered pairs.

Two systems take that test. A *model-free* character n -gram classifier over the occurrence’s context, with every instance of the variable’s name masked, and a fixed feature family chosen in advance. And a linear probe on the residual stream of Qwen2.5-Coder-1.5B [Hui et al., 2024], mean-pooled over the occurrence’s tokens, layer selected on source-language validation data. Intervals resample whole problems, never occurrences.

The surface classifier is strongly language-bound: 0.929 between JavaScript and PHP, 0.804 to or from Python, a boundary effect of -0.126 (95% CI $[-0.160, -0.092]$). The probe barely moves: -0.012 . And the standard defence holds: an untrained network of the same architecture and pooling reaches only 0.575 overall against the probe’s 0.844, a margin of $+0.269$. Every control in the usual battery says the probe is reading a learned, language-independent role representation.

3 The control the battery was missing

The asymmetry the setup hides is simple: the surface classifier is forbidden to read the variable’s name, and the probe is not. Parallel implementations of one problem tend to name their variables alike, so the name is exactly the kind of feature that transfers across languages for reasons that have nothing to do with the model.

54 The intervention is to change what the probe reads, not what the model sees. We pool the 16 tokens
55 either side of the occurrence and exclude its own, so the forward pass is byte-identical and only
56 the read-out moves. Masking the name in the source would instead change the input distribution,
57 entangling two effects.

58 Three things happen at once (Table 1, Figure 1). The probe falls from 0.844 to 0.589. Its margin
59 over a *matched* untrained floor collapses from +0.269 to +0.105: 61% of what the probe knew was
60 the name. And its boundary effect grows from -0.012 to -0.061 , against -0.064 for the untrained
61 network at the same pooling. With the name excluded, the trained probe is no flatter across languages
62 than an untrained one. The paired version holds per item: over the same occurrences, the surface
63 classifier beats the context-pooled probe in 14 of 18 cells with a problem-clustered interval excluding
64 zero.

65 **Why the untrained control missed it.** The untrained network is a control for *learning*, not for
66 *leakage*. It reads the same input as the trained probe, name included, and random projections of a
67 shared name support the same classifier on both sides of a language boundary. So the floor inherits
68 the confound: span-pooled, the untrained network’s boundary effect is -0.022 , nearly as flat as the
69 trained probe’s -0.012 . A control that shares the probe’s input cannot detect a property of the input.

70 **What caught it.** The model-free classifier, precisely because it was built with the name masked,
71 was the one condition in the design whose input did not contain the label’s correlate. Its strong
72 language-dependence is what made the probe’s flatness look remarkable, and the contradiction
73 between the two is what prompted the pooling intervention. A comparator that cannot see the suspect
74 feature is a different kind of control from a floor, and this case needed both.

75 4 The negative result

76 With the identifier excluded, what remains is a margin of +0.105 over the matched floor, consistent
77 in both language groups, so the model does encode something about roles beyond the name. But we
78 find no evidence that it is language-independent: its boundary effect is indistinguishable from an
79 untrained network’s. At this scale and in this setting, the language-independent role representation
80 the span-pooled probe appeared to demonstrate does not exist as measured.

81 We think the negative result is worth having on record because the positive version is easy to produce
82 and easy to believe. It requires only the standard methodology, it passes the standard controls, and
83 prior work in this area reports cross-language probe transfer with untrained floors and without a
84 model-free comparator [Hernández López et al., 2026, Utpala et al., 2024, Kargaran et al., 2025].
85 Our results do not show that those findings are wrong, but they do show that the design cannot rule
86 out this confound, and that in at least one natural setting the confound accounts for most of the effect.

87 5 Practical recommendations

88 **Vary the read-out, not only the model.** An untrained floor holds the probe’s input fixed and varies
89 the weights. Add at least one condition that varies what the probe may read: pool excluding the token
90 span whose label is being predicted, or compare against a model-free classifier denied the suspect
91 feature. The two catch different confounds.

92 **Match the floor to the read-out.** A context-pooled probe read against a span-pooled floor mixes two
93 chance levels. Whatever restriction the probe gets, its floor gets too.

94 **Treat input-label correlates as first-class suspects.** In parallel corpora, identifiers; in translation
95 corpora, named entities and numbers; in paired documents, any token the pairing preserves. If the
96 span being pooled contains a feature that survives the manipulation under study, the probe’s apparent
97 invariance to that manipulation is uninformative until the feature is excluded.

98 **Limitations.** One model at 1.5B, three languages, three roles, and a boundary instantiated by
99 a single language. Context pooling excludes the name but not everything correlated with it, so
100 +0.105 is an upper bound on name-free role information. The surface effect is concentrated in one
101 role (iterator, -0.225 alone against -0.059 for the other two), whose labels come from a different

102 extractor path across languages. And this remains a probing study: it bounds what is decodable, not
103 what the model uses.

104 References

105 Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational*
106 *Linguistics*, 48(1):207–219, 2022. doi: 10.1162/coli_a_00422.

107 José Antonio Hernández López, Martin Weyssow, Jesús Sánchez Cuadrado, and Houari Sahraoui.
108 Syntactic multilingual probing of pre-trained language models of code. *Journal of Systems and*
109 *Software*, 2026. doi: 10.1016/j.jss.2025.112604.

110 John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Proceedings*
111 *of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th*
112 *International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2733–
113 2743, Hong Kong, China, 2019. Association for Computational Linguistics. doi: 10.18653/v1/
114 D19-1275.

115 Binyuan Hui, Jian Yang, Zeyu Cui, Jiayi Yang, Dayiheng Liu, Lei Zhang, Tianyu Liu, Jiajun Zhang,
116 Bowen Yu, Keming Lu, Kai Dang, Yang Fan, Yichang Zhang, An Yang, Rui Men, Fei Huang,
117 Bo Zheng, Yibo Miao, Shanghaoran Quan, Yunlong Feng, Xingzhang Ren, Xuancheng Ren,
118 Jingren Zhou, and Junyang Lin. Qwen2.5-coder technical report. *arXiv preprint arXiv:2409.12186*,
119 2024. doi: 10.48550/arXiv.2409.12186.

120 Amir Hossein Kargaran, Yihong Liu, François Yvon, and Hinrich Schuetze. How programming
121 concepts and neurons are shared in code language models. In *Findings of the Association for*
122 *Computational Linguistics: ACL 2025*, 2025. doi: 10.18653/v1/2025.findings-acl.1379.

123 Jorma Sajaniemi. An empirical analysis of roles of variables in novice-level procedural programs. In
124 *Proceedings IEEE 2002 Symposia on Human Centric Computing Languages and Environments*,
125 pages 37–39, Arlington, VA, USA, 2002. IEEE Computer Society. doi: 10.1109/HCC.2002.
126 1046340.

127 Saiteja Utpala, Alex Gu, and Pin-Yu Chen. Language agnostic code embeddings. In *Proceedings*
128 *of the 2024 Conference of the North American Chapter of the Association for Computational*
129 *Linguistics: Human Language Technologies (Volume 1: Long Papers)*, 2024. doi: 10.18653/v1/
130 2024.naacl-long.38.

131 Elena Voita and Ivan Titov. Information-theoretic probing with minimum description length. In
132 *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*
133 *(EMNLP)*, 2020. doi: 10.18653/v1/2020.emnlp-main.14.

134 Kelly W. Zhang and Samuel R. Bowman. Language modeling teaches you more than translation
135 does: Lessons learned through auxiliary syntactic task analysis. In *Proceedings of the 2018*
136 *EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages
137 359–361, Brussels, Belgium, November 2018. Association for Computational Linguistics. doi:
138 10.18653/v1/W18-5448.

139 Ming Zhu, Aneesh Jain, Karthik Suresh, Roshan Ravindran, Sindhu Tipirneni, and Chandan K.
140 Reddy. XLCOST: A benchmark dataset for cross-lingual code intelligence. *arXiv preprint*
141 *arXiv:2206.08474*, 2022. doi: 10.48550/arXiv.2206.08474.

Table 2: Shared problem identifiers between every pair of XLCOST languages, train split. Matched cross-lingual transfer needs a pair to share enough problems to hold out a test fold; only the three bold cells clear 500. No pair outside Python/JavaScript/PHP does, including same-family pairs.

	C++	C#	Java	JavaScript	PHP	Python
C++	—	115	122	75	17	74
C#	115	—	175	75	14	62
Java	122	175	—	85	11	66
JavaScript	75	75	85	—	1,529	2,953
PHP	17	14	11	1,529	—	1,145
Python	74	62	66	2,953	1,145	—

Table 3: All eighteen transfer cells for Qwen2.5-Coder-1.5B. *Surface* is the strongest masked-context n -gram feature; *name* uses the variable’s name alone; *probe* is the residual-stream probe. Majority and shuffled-label controls sit near chance throughout, and ρ is the macro-F1 movement from one test occurrence changing its prediction.

role	pair	n	surface	name	probe	major.	shuf.	ρ
accumulator	JavaScript→PHP	1,976	0.951	0.897	0.904	0.361	0.470	0.0005
accumulator	JavaScript→Python	3,341	0.813	0.873	0.875	0.415	0.473	0.0004
accumulator	PHP→JavaScript	1,903	0.954	0.897	0.918	0.372	0.490	0.0005
accumulator	PHP→Python	1,377	0.782	0.824	0.833	0.276	0.439	0.0008
accumulator	Python→JavaScript	3,294	0.785	0.874	0.918	0.370	0.504	0.0003
accumulator	Python→PHP	1,478	0.736	0.845	0.874	0.271	0.466	0.0007
index_key	JavaScript→PHP	1,976	0.940	0.834	0.897	0.384	0.478	0.0005
index_key	JavaScript→Python	3,341	0.859	0.787	0.883	0.330	0.497	0.0003
index_key	PHP→JavaScript	1,903	0.964	0.838	0.880	0.376	0.516	0.0005
index_key	PHP→Python	1,377	0.839	0.782	0.854	0.387	0.515	0.0008
index_key	Python→JavaScript	3,294	0.871	0.822	0.915	0.314	0.470	0.0003
index_key	Python→PHP	1,478	0.858	0.838	0.874	0.419	0.459	0.0008
iterator	JavaScript→PHP	1,976	0.937	0.862	0.906	0.448	0.488	0.0008
iterator	JavaScript→Python	3,341	0.774	0.733	0.939	0.444	0.472	0.0005
iterator	PHP→JavaScript	1,903	0.965	0.873	0.851	0.446	0.478	0.0008
iterator	PHP→Python	1,377	0.576	0.788	0.865	0.428	0.462	0.0010
iterator	Python→JavaScript	3,294	0.790	0.769	0.918	0.466	0.484	0.0007
iterator	Python→PHP	1,478	0.741	0.788	0.902	0.475	0.438	0.0020

142 A Pairwise problem overlap in XLCOST

143 B Full transfer table

144 C Transfer mechanism

145 D Language geometry

146 **The typing axis is not privileged.** Probing the same activations under every alternative 4-versus-3
147 partition of the seven languages ($\binom{7}{4} = 35$ groupings minus the real one, so 34 controls) gives
148 best-layer macro-F1 0.9959 on average (min 0.9918, max 1.0000), against 1.0000 for the true
149 static/dynamic split; 5 of the 34 controls match or exceed it. Any grouping of these languages is
150 near-perfectly decodable, so the decodability of the typing split is evidence that language identity is
151 encoded, not that type discipline is. At layer 4, where $|r|$ peaks at 0.941, the remaining components
152 carry almost none of the label: PC2 $r = -0.014$, PC3 $r = +0.026$, PC4 $r = +0.171$.

153 E Causal interventions

154 **Scope.** Causal interventions were run for the *boolean* role over 5 languages (C++, Java, JavaScript,
155 PHP, Python), 3 models, and 3 intervention modes (ablate, patch, steer), giving 375 layer-wise

Table 4: Probe transfer per model, averaged within each group. The untrained network shares Qwen2.5-1.5B’s architecture and tokenizer with randomly initialised weights; it is the floor a representational claim must clear, and it is itself flat.

model	close pair	Python pairs	all	vs. untrained
Qwen2.5-1.5B	0.899	0.888	0.892	+0.316
Qwen2.5-Coder-1.5B	0.893	0.887	0.889	+0.314
Qwen2.5-1.5B (untrained)	0.590	0.568	0.575	—




Figure 2: Two candidate explanations for the transfer gap, over the six ordered pairs. Left: how much of the source classifier’s discriminative mass the target realises, which is essentially unrelated to transfer. Right: the share of that mass whose sign reverses between a source-fitted and a target-fitted classifier, which tracks it closely. The cues are present in both groups; what changes is where they point.

156 measurements. They are reported here rather than in the main text because they were run *within*
 157 languages and on a role outside the three this paper transfers, so they cannot support its central claim.
 158 Specificity is $|\text{clean} - \text{intervened}|$ divided by the same quantity for a random-position control: how
 159 much of the effect is specific to the variable’s own token positions rather than to editing anything at
 160 all. Two columns are given because the layer with the largest effect is generally not the layer with the
 161 best specificity, and quoting either alone misleads.

162 F Reproducibility

163 Every figure is regenerated from committed per-cell artifacts, and a regression test recomputes each
 164 from those files. All four conditions of Table 1 are backed by per-cell artifacts, with the one caveat
 165 the table’s footnote states. The probe is logistic regression on activations mean-pooled over either
 166 the occurrence’s span or the 16 tokens either side of it excluding the span; the surface classifier is
 167 a `char_wb` tf-idf vectoriser with $n \in [2, 4]$ and a pre-registered feature family, fitted on the source
 168 language only. Pooling and initialisation are recorded in each store’s identity, so conditions cannot be
 169 silently mixed downstream.

Table 5: Transfer mechanism for the iterator role. *Surviving* is the share of the source classifier’s discriminative mass realised in the target; *flip* is the share of that mass whose sign reverses. Surviving mass is uncorrelated with transfer; flipped mass predicts it almost exactly. Source-weighted variants are given for robustness.

pair	F1	surviving	flip	agree	flip (src-wt.)
PHP→JavaScript	0.965	0.735	0.072	+0.895	0.072
JavaScript→PHP	0.937	0.698	0.052	+0.945	0.068
Python→JavaScript	0.790	0.862	0.178	+0.695	0.211
JavaScript→Python	0.774	0.757	0.232	+0.811	0.241
Python→PHP	0.741	0.696	0.280	+0.528	0.367
PHP→Python	0.576	0.697	0.366	+0.476	0.346

Table 6: Masking an entire syntactic character class on both sides of the transfer. No class accounts for more than 0.021 of the 0.39 gap, so the realignment is distributed rather than carried by block delimiters or statement terminators.

pair	masked class	macro-F1	Δ
PHP→JavaScript	terminator	0.9496	-0.0159
PHP→JavaScript	brace	0.9645	-0.0010
PHP→JavaScript	bracket	0.9711	+0.0057
PHP→JavaScript	paren	0.9609	-0.0045
PHP→JavaScript	operator	0.9441	-0.0213
PHP→Python	terminator	0.5786	+0.0030
PHP→Python	brace	0.5750	-0.0006
PHP→Python	bracket	0.5733	-0.0023
PHP→Python	paren	0.5947	+0.0191
PHP→Python	operator	0.5801	+0.0045

Table 7: Point-biserial correlation between PC1 of last-token residual activations and a static/dynamic typing label, by layer, with PC1’s explained-variance ratio. The sign of r is arbitrary (PCA fixes the axis only up to reflection, and each layer is refit independently); magnitude is what carries meaning.

layer	$ r $	EVR	layer	$ r $	EVR
0	0.883	0.481	15	0.894	0.309
1	0.898	0.447	16	0.882	0.303
2	0.914	0.391	17	0.874	0.296
3	0.939	0.373	18	0.873	0.289
4	0.941	0.370	19	0.873	0.299
5	0.862	0.369	20	0.903	0.291
6	0.828	0.361	21	0.886	0.276
7	0.844	0.342	22	0.859	0.265
8	0.896	0.339	23	0.842	0.252
9	0.919	0.333	24	0.812	0.242
10	0.880	0.327	25	0.774	0.253
11	0.893	0.302	26	0.694	0.255
12	0.917	0.292	27	0.732	0.237
13	0.904	0.314	28	0.576	0.242
14	0.898	0.309			

Table 8: Causal interventions on boolean-flag occurrences. *Peak-effect spec.* is specificity at the largest-effect layer; *best spec.* is the maximum over layers, with that layer named.

language	mode	model	peak-effect spec.	best spec. (layer)
C++	ablate	qwen2515b	4.5×	31.9× (L4)
C++	ablate	qwen25coder15b	4.4×	76.2× (L20)
C++	ablate	starcoder27b	4.3×	36.5× (L4)
C++	patch	qwen2515b	3.6×	7.9× (L20)
C++	patch	qwen25coder15b	3.6×	6.9× (L8)
C++	patch	starcoder27b	8.1×	10.4× (L16)
C++	steer	qwen2515b	45.0×	169.8× (L4)
C++	steer	qwen25coder15b	598.5×	598.5× (L24)
C++	steer	starcoder27b	6.9×	341.1× (L4)
Java	ablate	qwen2515b	28.5×	86.0× (L12)
Java	ablate	qwen25coder15b	12.3×	62.4× (L12)
Java	ablate	starcoder27b	27.0×	27.0× (L4)
Java	patch	qwen2515b	6.1×	7.7× (L12)
Java	patch	qwen25coder15b	5.9×	7.7× (L12)
Java	patch	starcoder27b	7.5×	8.6× (L12)
Java	steer	qwen2515b	72.1×	849.0× (L8)
Java	steer	qwen25coder15b	71.9×	186.6× (L0)
Java	steer	starcoder27b	252.0×	252.0× (L24)
JavaScript	ablate	qwen2515b	26.7×	26.7× (L16)
JavaScript	ablate	qwen25coder15b	28.1×	116.3× (L24)
JavaScript	ablate	starcoder27b	27.3×	27.3× (L4)
JavaScript	patch	qwen2515b	6.9×	6.9× (L8)
JavaScript	patch	qwen25coder15b	6.6×	7.4× (L12)
JavaScript	patch	starcoder27b	8.1×	8.2× (L16)
JavaScript	steer	qwen2515b	24.5×	76.4× (L8)
JavaScript	steer	qwen25coder15b	65.2×	149.5× (L4)
JavaScript	steer	starcoder27b	11.2×	60.8× (L28)
PHP	ablate	qwen2515b	15.0×	102.4× (L20)
PHP	ablate	qwen25coder15b	26.5×	79.2× (L20)
PHP	ablate	starcoder27b	19.4×	19.7× (L8)
PHP	patch	qwen2515b	7.2×	7.2× (L4)
PHP	patch	qwen25coder15b	4.7×	7.4× (L0)
PHP	patch	starcoder27b	6.3×	9.3× (L16)
PHP	steer	qwen2515b	190.3×	190.3× (L24)
PHP	steer	qwen25coder15b	11.5×	134.0× (L8)
PHP	steer	starcoder27b	8.7×	100.1× (L4)
Python	ablate	qwen2515b	5.8×	36.7× (L24)
Python	ablate	qwen25coder15b	6.8×	33.1× (L16)
Python	ablate	starcoder27b	26.5×	26.5× (L4)
Python	patch	qwen2515b	6.0×	7.2× (L16)
Python	patch	qwen25coder15b	6.7×	7.8× (L16)
Python	patch	starcoder27b	7.9×	9.7× (L12)
Python	steer	qwen2515b	209.1×	827.9× (L16)
Python	steer	qwen25coder15b	1025.8×	1025.8× (L24)
Python	steer	starcoder27b	8.6×	55.5× (L4)

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”,**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and Section 1 state the transfer figures, the untrained floor, and the mechanism result, and each is the number reported in Sections 4 and 5. Claims are scoped to three languages, three roles, and two trained models throughout.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The Limitations paragraph of Section 6 covers the two-point typological axis, the single model family and scale, the extractor-definition difference affecting the iterator role, the scope of the mechanism analysis, and the absence of causal evidence on the cross-lingual cells.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper reports no theoretical results.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 2 gives the corpus, the matching procedure, the occurrence cap, the split policy, the baseline feature families, and the probe’s layer-selection rule. Appendix F names the scripts, the model identifiers, and the classifier settings.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All result tables are committed as CSVs under `results/lp4fm/`, and the analysis scripts are released. Appendix F gives the commands. An anonymised mirror is provided for the review period.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.

- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 2 specifies the cap, the frozen hash split, and the source-validation layer selection. Appendix F gives the n -gram vectoriser configuration and the logistic-regression settings.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: We report resolution (ρ , the macro-F1 movement from one test occurrence changing its prediction) for every cell rather than confidence intervals, and state in Section 2 which comparisons clear it and which do not. Per-cell cluster-bootstrap intervals are computed by the released code but are not reported in this version; the mechanism correlations rest on six ordered pairs and are reported with intervals in Section 5.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix F reports the accelerator used and the wall-clock cost of activation extraction, which dominates. The surface baselines and the mechanism analyses are CPU-only.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work analyses publicly released models on a public benchmark of competitive-programming solutions. It involves no human subjects, no personally identifying data, and no deployment.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: The work is an analysis of how existing code models represent variable roles. We identify no societal impact specific to it beyond that of the public models and benchmark it studies.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.

- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases no data or models. It analyses publicly available ones.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: XLCoST is cited and its licence respected. The two Qwen model releases are cited by name and identifier in Section 2 and Appendix F.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The released analysis scripts and result files are documented in Appendix F, including how each table is regenerated from the committed CSVs.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper involves no crowdsourcing and no human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper involves no human subjects, so no IRB review was required.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [No]

533 Justification: Large language models are the object of study, not a component of the method.
534 They were used for writing and editing assistance only, which the question exempts from
535 declaration.

536 Guidelines:

- 537 • The answer [N/A] means that the core method development in this research does not
538 involve LLMs as any important, original, or non-standard components.
- 539 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
540 be described.