

---

# Same Score, Different Evidence: Decodability, Surface Sufficiency, and Causal Relevance in Code Models

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 Interpretability experiments must distinguish model mechanisms from measure-  
2 ment artifacts. We present a case-study-backed reporting matrix for linear probes  
3 over five identifier properties and roles in code models. The study includes three  
4 models and collectively covers all seven XLCOST programming languages without  
5 forming a complete factorial design. The matrix binds each claim to an estimand,  
6 comparator, matching unit, uncertainty statement, outcome, and falsifier. Two  
7 matched cases motivate it. On Python, boolean occurrence type is decodable at  
8 0.981–0.988 macro-F1, while a masked source-line classifier reaches 0.983 with-  
9 out a language model. Paired problem-clustered deltas are  $-0.003$  to  $+0.002$ , with  
10 intervals excluding probe advantages above 0.014–0.017, although the sign of that  
11 difference varies by language and by comparator. A full-residual patch at a class or  
12 structure query-name site also fails its specified bidirectional gate in Qwen2.5-1.5B,  
13 recovering 0.009 to 0.020 of the matched behavioral gap. This bounds the effect at  
14 that site rather than showing global non-use. A third case exposes a failed control.  
15 Role-conditioned iterator renaming yields 0.998–0.999 F1 at embedding layer zero,  
16 consistent with a lexical code introduced by the intervention rather than semantic  
17 robustness. Decodability, capacity control, cue sensitivity, surface sufficiency, and  
18 site-state causal relevance therefore require separate, outcome-aware records.

## 19 1 Introduction

20 Linear probes are often used to ask what information a representation contains. A successful  
21 probe establishes decodability under a particular readout. It does not by itself establish that the  
22 model computed an abstract feature, that the feature is unavailable in surface form, or that the  
23 model uses it behaviorally [Belinkov, 2022, Voita and Titov, 2020]. Control tasks constrain probe  
24 capacity [Hewitt and Liang, 2019], and targeted studies show that probe accuracy need not entail task  
25 relevance [Ravichander et al., 2021, Elazar et al., 2021]. Capacity is only one source of ambiguity.  
26 Input leakage, label construction, sample resolution, and an unmatched intervention can each preserve  
27 high accuracy while changing the scientific claim.

28 This distinction is central to interpretability as a science. A measurement is useful when its target  
29 is explicit, alternative explanations are tested, and contrary outcomes change the conclusion. We  
30 use identifier probing as a comparative case study. Some targets are semantic variable roles, such  
31 as accumulating values or indexing a container [Sajaniemi, 2002]. Others are occurrence types or  
32 declared type identifiers. Names, syntax, and context can all predict these targets, which makes probe  
33 accuracy underidentified. Prior code-probing studies, including the 15-task INSPECT framework,  
34 map decodable syntactic and semantic information across layers and models [Karmakar and Robbes,  
35 2021, Troshin and Chirkova, 2022, Karmakar and Robbes, 2024]. We do not propose another probe,

Table 1: Target coverage. No row should be read as a complete three-model by seven-language experiment. Table 3 gives the outcome-aware evidence records.

| Target             | Models       | Language coverage                         |
|--------------------|--------------|-------------------------------------------|
| Accumulator        | Qwen2.5-1.5B | Python plus Java, C++, C, C#              |
| Index or key       | Qwen2.5-1.5B | Python, Java, JavaScript, C++, C, C#      |
| Iterator           | All three    | Python, Java, JavaScript, PHP, C++, C, C# |
| Boolean            | All three    | Python, Java, JavaScript, PHP             |
| Class or structure | All three    | Python, C++, JavaScript, C                |

control, or task taxonomy. Existing work catalogs probing limitations and shows that activation-patching conclusions depend on metric and intervention choices [Belinkov, 2022, Zhang and Nanda, 2023, Heimersheim and Nanda, 2024]. Subspace interventions can also create an illusion of localization [Makelov et al., 2023], which motivates calls for falsifiable interpretability designs [Leavitt and Morcos, 2020] and for naming the causal mediator a claim is about [Mueller et al., 2024]. Reporting standards in other empirical sciences answer a related need by binding a published claim to its estimand, comparator, and uncertainty statement [Schulz et al., 2010]. Our distinct contribution is an outcome-aware reporting unit that binds each claim to its estimand, comparator, matching unit, uncertainty, observed outcome, and falsifying next test. Unlike a control checklist, the record preserves failed and confounded comparisons and states how they change the licensed claim.

Our evidence covers five roles, three models, and the seven-language XLCOST corpus [Zhu et al., 2022]. The experiment is intentionally reported as an uneven evidence matrix rather than a full factorial study. Qwen2.5-1.5B, Qwen2.5-Coder-1.5B, and StarCoder2-7B all appear in the controlled boolean, class or structure, and iterator experiments [Qwen et al., 2024, Hui et al., 2024, Lozhkov et al., 2024]. Earlier accumulator and index runs use one model and a weaker split protocol. Iterator runs cover all three models and all seven languages. This is a methodological case study of evidential boundaries, not a validation of the matrix as a universal framework. Its heterogeneity separates replicated findings from leads that still require confirmation.

The central result is that high probe F1 produces different evidential outcomes. The paired boolean comparison supports a bounded negative. The class or structure intervention fails its gate at one site and layer. Role-conditioned iterator renaming produces near-perfect embedding-layer scores, revealing that a nominal control can introduce a new shortcut. Legacy accumulator and index effects remain motivating observations. A single accuracy number, or a checkmark indicating that a control ran, cannot identify among these explanations.

## 2 Study design and evidence matrix

**Targets, models, and languages.** We treat five identifier properties and roles as related case studies, not measurement-equivalent instances of one construct. Accumulators update within loops. Indices or keys appear as counters or subscripts. Iterators are loop-bound names and their uses. Boolean labels distinguish assignment, conditional, loop, and return occurrences. Class or structure labels mark declared type names and references. Binary token probes label all other tokens as negative. Language-specific AST and regular-expression extractors lack a common multilingual gold audit. Protocols also differ. Accumulator and index use legacy token splits and test-maximized layers. Boolean mean-pools occurrences and uses problem groups. Iterator and class or structure use token labels, problem groups, validation-selected layers, and five seeds. XLCOST contains C, C++, C#, Java, JavaScript, PHP, and Python. Table 1 states coverage. Scores across rows do not estimate one common task.

**Four comparison questions.** Renaming asks whether lexical identity carries the prediction. Random-noun, single-character, and numeric strategies are role-agnostic textual substitutions. The misleading strategy is not role-agnostic. Target identifiers receive counter-role names while other identifiers receive role-looking names. The vocabulary therefore encodes the label. The all-same strategy is also non-bijective. Because every condition refits its probe and layer, neither strategy tests semantic robustness or a fixed invariant direction. We retain them as diagnostics of a failed intervention. A model-free character classifier with the target name masked asks whether local

Table 2: Case outcomes after interpreting the available comparisons. Values are macro-F1. A diagnostic can reveal a failed control without supporting a representation claim.

| Case                | Probe        | Comparison outcome                                | Strongest supported inference                                 |
|---------------------|--------------|---------------------------------------------------|---------------------------------------------------------------|
| Legacy accum./index | 0.893, 0.959 | $\Delta = +0.079, -0.272$                         | Exploratory cue response under a weaker protocol              |
| Iterator            | 0.976–0.987  | role-conditioned rename reaches 0.998–0.999 at L0 | The intervention encodes a lexical shortcut                   |
| Boolean             | 0.981–0.988  | paired $\Delta \in [-0.003, +0.002]$              | Probe advantage over masked local source excluded above 0.017 |
| Class/site          | 0.979–0.982  | patch recovery 0.009, 0.020                       | Bounded site-state effect at the tested site and layer        |

source text already determines the label. Randomized-label controls ask whether the probe family can fit a specified null. Boolean assigns a random label per variable name. Iterator and class or structure permute labels within programs. Full-residual patching asks whether replacing the complete site state moves a behavioral logit readout [Meng et al., 2022]. Its interpretation depends on the corruption, metric, and readout, so these choices belong in the evidence record [Zhang and Nanda, 2023, Heimersheim and Nanda, 2024].

These comparisons are not interchangeable. Role-agnostic refit renaming can measure recoverability after a cue changes, but role-conditioned renaming can create a cue. Positive selectivity can reject its specified permutation task while leaving a surface shortcut intact. A model-free comparator can establish surface sufficiency but cannot show how the model uses the feature. Full-residual patching tests site-state relevance, not use of the decoded probe direction, and only at its specified layer, span, direction, and readout.

**Evidence grades.** Boolean, iterator, and class or structure experiments use problem groups, validation-selected layers, five seeds, and BCa or clustered uncertainty. Iterator additionally reports a within-program label-permutation control and a name-only comparator, but their difference is not a paired clustered interval. Accumulator and index use token splits, test-maximized layers, and whole-word renaming. We retain their opposite effects as legacy hypotheses, not population estimates. A DeepSeek index run is excluded because a tokenizer offset error invalidates its activation alignment.

### 3 The same probe score supports different claims

#### 3.1 Renaming diagnoses an intervention confound

Legacy Qwen2.5-1.5B runs first motivated the comparison. The index probe falls from 0.959 to 0.687 under misleading names, a change of  $-0.272$ , while the accumulator probe rises from 0.893 to 0.972, a change of  $+0.079$ . Token-level splitting, test-maximized layers, role-conditioned replacements, and refitting prevent a robustness interpretation. These are exploratory cue responses rather than a confirmed cross-role dissociation.

The modern iterator rerun reveals the intervention problem directly. Across the three models, baseline F1 is 0.976, 0.978, and 0.987. The widest BCa 95% interval is  $[0.964, 0.982]$ , and permutation selectivity is  $+0.526$  to  $+0.552$ . Role-conditioned misleading names change F1 by  $+0.023$ ,  $+0.020$ , and  $+0.012$ , but all three probes select embedding layer zero and reach 0.998–0.999. The name pools encode target membership, so this result is compatible with relearning an inverted lexical rule. It does not support semantic robustness. Single-character and numeric names reduce F1 by 0.046–0.072. A name-only classifier on the same 427 Python programs reaches 0.918, leaving an unpaired gap of 0.058–0.069. Python-trained transfer ranges are 0.934–0.981, 0.898–0.923, and 0.887–0.938. Because target evaluation is not separated by one global cross-language problem partition, these scores establish descriptive transfer rather than language abstraction. Exploratory iterator patching appears in the appendix and is not the class or structure causal gate.

### 3.2 A model-free baseline bounds the boolean result

Boolean occurrence type provides the clearest bounded negative result. On a frozen Python sample, the three models reach 0.982, 0.981, and 0.988 macro-F1. Selectivity ranges from +0.272 to +0.288. These values reject the specific hypothesis that an equally expressive probe predicts labels randomly assigned per variable name as well as it predicts the real usage labels. Yet a character classifier on the enclosing source line, with the variable name masked, reaches 0.983. The probe-minus-surface differences are  $-0.001$ ,  $-0.002$ , and  $+0.005$ .

The original exports retained no aligned predictions, so we reran both predictors on the frozen sample and used aligned occurrence identifiers. Pooling the five seed test folds gives 2,067 seed-level predictions for 1,410 unique occurrences in 915 of the 1,301 source problems, and predictor overlap is complete on every fold. The committed artifact retains the aggregate paired summary, not those prediction-level inputs. The probe-minus-surface estimates and problem-clustered 95% intervals are  $-0.002$   $[-0.017, 0.014]$ ,  $-0.003$   $[-0.037, 0.014]$ , and  $+0.002$   $[-0.016, 0.017]$ . Every interval covers zero and excludes a probe advantage larger than 0.014–0.017 macro-F1. The masked-line comparator was the best of the reported surface feature families on this frozen sample, so the bound is conditional on that non-nested selection. That selection matters. Against the next-best Python comparator, the masked enclosing statement at 0.970, the differences are  $+0.010$  to  $+0.017$ . Across the four committed languages, probe minus best baseline ranges from  $-0.004$  to  $+0.021$ , the largest value being PHP with StarCoder2-7B, and line masking is degenerate outside Python because detokenized XLCoST programs there occupy a single line. Five refit renaming conditions change F1 by at most 0.013, but the role-conditioned and non-bijective strategies prevent a robustness claim. The supported conclusion is narrower. On Python, the probe does not outperform the selected masked local-syntax comparator by more than 0.014–0.017. The other languages have unpaired differences only, so they carry no bound.

### 3.3 A site-state intervention fails its specified gate

The class or structure experiments and the boolean experiments use different controls. Across all three models, baseline F1 is 0.979–0.982 and within-program permutation selectivity is  $+0.503$ – $+0.553$ . Misleading class-style names change F1 by  $-0.004$ ,  $-0.004$ , and  $+0.004$ , but this role-conditioned vocabulary prevents a robustness interpretation. Python-trained probes transfer to C++, JavaScript, and C at 0.885–0.980. These results support decodability and performance above the specified within-program permutation task. They do not establish abstraction or a fixed rename-invariant direction.

We then test site-state relevance in Qwen2.5-1.5B by patching the full residual state. The experiment contains 288 matched class-versus-function prompts in 48 template clusters. Let  $D = \ell(\text{True}) - \ell(\text{False})$ . Denoising patches the class-source residual into a function prompt and defines  $e_d = D_{\text{patched function}} - D_{\text{function}}$ . Noising reverses the patch and defines  $e_n = D_{\text{class}} - D_{\text{patched class}}$ . Recovery is the ratio of the mean signed effect to the mean matched gap  $D_{\text{class}} - D_{\text{function}}$ . Percentile intervals resample template clusters. The gate requires positive, control-separated effects, recovery of at least 0.05, and mean effects of at least 0.10 in both directions. These thresholds are design choices rather than calibrated behavioral minima.

Both prompt classes prefer True, which limits the behavioral readout’s effective decision range. Their mean logit differences are +2.75 for class prompts and +1.78 for function prompts, although the matched class-minus-function gap is +0.97 and has the expected sign for 98.6% of pairs. The gate therefore tests whether patching transfers this relative distinction, not whether it flips a well-calibrated binary decision.

Denoising produces a mean effect of 0.0087 with a 95% interval  $[-0.0007, 0.0188]$  and recovery 0.0089. Noising produces 0.0199  $[0.0109, 0.0296]$  with recovery 0.0204. The full-residual intervention fails the specified gate. The protocol and results entered the repository together, so we do not claim auditable preregistration. Recovery divides by the matched gap, so a compressed decision range does not explain a ratio of 0.009 to 0.020. We read this as a bounded negative at the tested site, layer, and span rather than evidence that class information is absent or unused. The intervals exclude mean effects above 0.0188 and 0.0296 logit-difference units in the tested directions. Decodability and selectivity therefore do not establish a substantive site-state contribution under this readout.

Table 3: Outcome-aware evidence records. “Unpaired” means that uncertainty on the difference is unavailable. Each next test can change the stated claim.

| Case         | Comparison and outcome                            | Matching and uncertainty                      | Claim and falsifier                                                                                   |
|--------------|---------------------------------------------------|-----------------------------------------------|-------------------------------------------------------------------------------------------------------|
| Legacy roles | Refit accumulator +0.079, index −0.272            | Token split, test-selected layer, no interval | Exploratory only. Rerun grouped with label-blind renames                                              |
| Iterator     | Role-conditioned rename reaches L0 at 0.998–0.999 | Problem split and marginal BCa intervals      | Control is confounded. Use bijective role-agnostic renames                                            |
| Boolean      | Probe minus masked line is −0.003 to +0.002       | 915 clusters with paired 95% intervals        | Not supported on Python. No selected-baseline advantage above 0.017. Nest selection and add languages |
| Class/site   | Bidirectional residual-patch gate fails           | 48 template clusters with 95% intervals       | Bounded effect at one site and layer. Test other spans and layers                                     |

## 169 4 An outcome-aware reporting matrix

170 The contribution is a reporting obligation, not a new control or evidence ladder. Each record states the  
171 target and prediction unit, claim and estimand, comparator and matching rule, uncertainty, outcome,  
172 and falsifying next test. We distinguish five claim types. Decodability exceeds a label baseline.  
173 Capacity control exceeds a specified randomized-label task. Cue sensitivity uses an intervention that  
174 does not itself encode the label. Surface sufficiency compares a model-free predictor on a matched  
175 population. Site-state causal relevance requires a gate-specified intervention and behavioral readout.  
176 No claim entails another.

177 These heterogeneous cases demonstrate how the record changes interpretation. They do not validate  
178 the matrix as universal. A claim of cross-language abstraction remains open until one extractor,  
179 global problem partition, and matched control battery cover a complete language grid. Replication  
180 packages should publish the record beside metric tables and state the smallest effect the design can  
181 exclude.

## 182 5 Conclusion

183 High probe accuracy is not a common unit of interpretability evidence. In these cases, matched  
184 surface prediction supports a bounded negative, role-conditioned renaming exposes a shortcut, and a  
185 specified patching gate fails at one site and layer. Publishing each claim as an outcome-aware record  
186 lets null and confounded outcomes constrain conclusions rather than disappear behind a control  
187 checkmark.

## 188 6 Limitations

189 The study collectively covers three models, five roles, and seven languages, but not a complete  
190  $3 \times 5 \times 7$  design, and Table 1 gives the per-target coverage. The targets are heterogeneous, and  
191 their language-specific extractors lack a common gold audit. Role-conditioned and all-same renames  
192 can encode or alter labels. Legacy accumulator-index results lack problem groups and scope-aware  
193 renaming. Iterator has no paired probe-minus-surface interval, and its target evaluation does not  
194 enforce one global cross-language problem partition. The boolean bound is conditional on a surface  
195 comparator selected on the frozen sample, its sign varies across languages, and its prediction-level  
196 paired inputs are not retained. Full-residual patching covers one model, layer, site, and readout. The  
197 gate and stopping rule were not registered in advance. These cases motivate the reporting matrix  
198 rather than validate it.

## References

- Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219, 2022. doi: 10.1162/coli\_a\_00422.
- Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav Goldberg. Amnesic probing: Behavioral explanation with amnesic counterfactuals. *Transactions of the Association for Computational Linguistics*, 9:160–175, 2021. doi: 10.1162/tacl\_a\_00359.
- Stefan Heimersheim and Neel Nanda. How to use and interpret activation patching. *arXiv preprint arXiv:2404.15255*, 2024. doi: 10.48550/arXiv.2404.15255.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2733–2743, Hong Kong, China, 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1275.
- Binyuan Hui, Jian Yang, Zeyu Cui, Jiaxi Yang, Dayiheng Liu, Lei Zhang, Tianyu Liu, Jiajun Zhang, Bowen Yu, Keming Lu, Kai Dang, Yang Fan, Yichang Zhang, An Yang, Rui Men, Fei Huang, Bo Zheng, Yibo Miao, Shanghaoran Quan, Yunlong Feng, Xingzhang Ren, Xuancheng Ren, Jingren Zhou, and Junyang Lin. Qwen2.5-coder technical report. *arXiv preprint arXiv:2409.12186*, 2024. doi: 10.48550/arXiv.2409.12186.
- Anjan Karmakar and Romain Robbes. What do pre-trained code models know about code? In *2021 36th IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 1332–1336, Melbourne, Australia, 2021. IEEE. doi: 10.1109/ASE51524.2021.9678927.
- Anjan Karmakar and Romain Robbes. INSPECT: Intrinsic and systematic probing evaluation for code transformers. *IEEE Transactions on Software Engineering*, 50(2):220–238, 2024. doi: 10.1109/TSE.2023.3341624.
- Matthew L. Leavitt and Ari Morcos. Towards falsifiable interpretability research. *arXiv preprint arXiv:2010.12016*, 2020. doi: 10.48550/arXiv.2010.12016.
- Anton Lozhkov, Raymond Li, Loubna Ben Allal, Federico Cassano, Joel Lamy-Poirier, Nouamane Tazi, Ao Tang, Dmytro Pykhtar, Jiawei Liu, Yuxiang Wei, Tianyang Liu, Max Tian, Denis Kocetkov, Arthur Zucker, Younes Belkada, Zijian Wang, Qian Liu, Dmitry Abulkhanov, Indraneil Paul, Zhuang Li, Wen-Ding Li, Megan Risdal, Jia Li, Jian Zhu, Terry Yue Zhuo, Evgenii Zheltonozhskii, Nii Osa Osa Dade, Wenhao Yu, Lucas Krauss, Naman Jain, Yixuan Su, Xuanli He, Manan Dey, Edoardo Abati, Yekun Chai, Niklas Muennighoff, Xiangru Tang, Muhtasham Oblokulov, Christopher Akiki, Marc Marone, Chenghao Mou, Mayank Mishra, Alex Gu, Binyuan Hui, Tri Dao, Armel Zebaze, Olivier Dehaene, Nicolas Patry, Canwen Xu, Julian McAuley, Han Hu, Torsten Scholak, Sebastien Paquet, Jennifer Robinson, Carolyn Jane Anderson, Nicolas Chapados, Mostofa Patwary, Nima Tajbakhsh, Yacine Jernite, Carlos Munoz Ferrandis, Lingming Zhang, Sean Hughes, Thomas Wolf, Arjun Guha, Leandro von Werra, and Harm de Vries. StarCoder 2 and The Stack v2: The next generation. *arXiv preprint arXiv:2402.19173*, 2024. doi: 10.48550/arXiv.2402.19173.
- Aleksandar Makelov, Georg Lange, and Neel Nanda. Is this the subspace you are looking for? an interpretability illusion for subspace activation patching. *arXiv preprint arXiv:2311.17030*, 2023. doi: 10.48550/arXiv.2311.17030.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. *Advances in Neural Information Processing Systems*, 35, 2022. doi: 10.52202/068431-1262.
- Aaron Mueller, Jannik Brinkmann, Millicent Li, Samuel Marks, Koyena Pal, Nikhil Prakash, Can Rager, Aruna Sankaranarayanan, Arnab Sen Sharma, Jiuding Sun, Eric Todd, David Bau, and Yonatan Belinkov. The quest for the right mediator: Surveying mechanistic interpretability through the lens of causal mediation analysis. *arXiv preprint arXiv:2408.01416*, 2024. doi: 10.48550/arXiv.2408.01416.

- Qwen, An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tianyi Tang, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.
- Abhilasha Ravichander, Yonatan Belinkov, and Eduard Hovy. Probing the probing paradigm: Does probing accuracy entail task relevance? In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 3363–3377, Online, 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.eacl-main.295.
- Jorma Sajaniemi. An empirical analysis of roles of variables in novice-level procedural programs. In *Proceedings IEEE 2002 Symposia on Human Centric Computing Languages and Environments*, pages 37–39, Arlington, VA, USA, 2002. IEEE Computer Society. doi: 10.1109/HCC.2002.1046340.
- Kenneth F. Schulz, Douglas G. Altman, and David Moher. CONSORT 2010 statement: Updated guidelines for reporting parallel group randomised trials. *BMJ*, 340:c332, 2010. doi: 10.1136/bmj.c332.
- Sergey Troshin and Nadezhda Chirkova. Probing pretrained models of source codes. In *Proceedings of the Fifth BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 371–383, Abu Dhabi, United Arab Emirates (Hybrid), December 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.blackboxnlp-1.31. URL <https://aclanthology.org/2022.blackboxnlp-1.31/>.
- Elena Voita and Ivan Titov. Information-theoretic probing with minimum description length. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 183–196, 2020. doi: 10.18653/v1/2020.emnlp-main.14.
- Fred Zhang and Neel Nanda. Towards best practices of activation patching in language models: Metrics and methods. *arXiv preprint arXiv:2309.16042*, 2023. doi: 10.48550/arXiv.2309.16042.
- Ming Zhu, Aneesh Jain, Karthik Suresh, Roshan Ravindran, Sindhu Tipirneni, and Chandan K. Reddy. XLCoS: A benchmark dataset for cross-lingual code intelligence. *arXiv preprint arXiv:2206.08474*, 2022. doi: 10.48550/arXiv.2206.08474.

## 279 A Per-role control results

Table 4: Complete boolean occurrence-type probe and model-free baseline results. The best baseline is selected from name-only, masked statement, masked line, masked window, covariate-only, and majority features.  $\rho$  is a descriptive one-instance score step, not an interval or equivalence test. Aligned probe/baseline predictions were not retained.

| language   | model              | problems | probe F1 | best baseline | difference | $\rho$ | selectivity |
|------------|--------------------|----------|----------|---------------|------------|--------|-------------|
| Java       | Qwen2.5-1.5B       | 1,427    | 0.982    | 0.980         | +0.002     | 0.0120 | +0.255      |
| Java       | Qwen2.5-Coder-1.5B | 1,427    | 0.976    | 0.980         | -0.004     | 0.0079 | +0.271      |
| Java       | StarCoder2-7B      | 1,427    | 0.990    | 0.980         | +0.010     | 0.0120 | +0.243      |
| JavaScript | Qwen2.5-1.5B       | 1,351    | 0.989    | 0.992         | -0.003     | 0.0170 | +0.132      |
| JavaScript | Qwen2.5-Coder-1.5B | 1,351    | 0.998    | 0.992         | +0.006     | 0.0170 | +0.150      |
| JavaScript | StarCoder2-7B      | 1,351    | 0.988    | 0.992         | -0.004     | 0.0170 | +0.120      |
| PHP        | Qwen2.5-1.5B       | 1,305    | 0.983    | 0.969         | +0.014     | 0.0080 | +0.214      |
| PHP        | Qwen2.5-Coder-1.5B | 1,305    | 0.975    | 0.969         | +0.006     | 0.0080 | +0.216      |
| PHP        | StarCoder2-7B      | 1,305    | 0.990    | 0.969         | +0.021     | 0.0080 | +0.240      |
| Python     | Qwen2.5-1.5B       | 1,301    | 0.982    | 0.983         | -0.001     | 0.0068 | +0.280      |
| Python     | Qwen2.5-Coder-1.5B | 1,301    | 0.981    | 0.983         | -0.002     | 0.0068 | +0.272      |
| Python     | StarCoder2-7B      | 1,301    | 0.988    | 0.983         | +0.005     | 0.0068 | +0.288      |

Table 5: Boolean Python identifier-renaming audit. C1–C5 are the five committed renaming conditions. Each condition refits a probe and reselects its layer, so deltas measure recoverability rather than fixed-direction invariance. Intervals are problem-clustered. Empty renaming fields for other languages indicate that those runs were not committed, not a null result.

| model              | C1     | C2     | C3     | C4     | C5     |
|--------------------|--------|--------|--------|--------|--------|
| Qwen2.5-1.5B       | +0.003 | -0.004 | +0.001 | -0.003 | +0.001 |
| Qwen2.5-Coder-1.5B | +0.001 | -0.010 | +0.006 | -0.004 | -0.006 |
| StarCoder2-7B      | -0.003 | -0.013 | -0.006 | -0.001 | -0.006 |

280 **Renaming diagnostic.** The misleading strategy uses role-conditioned vocabularies. Target and  
 281 other identifiers receive names from different pools, so a refit probe can recover the label lexically.  
 282 The all-same strategy is non-bijective and can alter label extraction. These rows diagnose intervention  
 283 confounds rather than semantic robustness.

Table 6: Complete model-free boolean baselines by language. Baselines are shared across the three neural models because they use the same frozen examples and surface features.

| language   | majority | name  | statement | line  | window | covariates |
|------------|----------|-------|-----------|-------|--------|------------|
| Java       | 0.311    | 0.522 | 0.980     | 0.501 | 0.662  | 0.311      |
| JavaScript | 0.227    | 0.373 | 0.992     | 0.369 | 0.552  | 0.226      |
| PHP        | 0.228    | 0.505 | 0.969     | 0.413 | 0.581  | 0.228      |
| Python     | 0.217    | 0.422 | 0.970     | 0.983 | 0.605  | 0.216      |

284 **Boolean probe figures.** The baseline panels retain the individual surface comparators hidden by  
 285 the best-baseline table. Renaming is available only for Python. The absence of corresponding figures  
 286 in other languages is a coverage hole.

Table 7: Complete class/structure perturbation results under the shared five-seed protocol. Each condition refits a probe and selects its own layer on validation data, so cross-condition deltas do not test a fixed direction.

| model              | strategy                | layer | F1    | 95% CI         | control F1 | selectivity | $\Delta$ |
|--------------------|-------------------------|-------|-------|----------------|------------|-------------|----------|
| Qwen2.5-1.5B       | baseline                | L18   | 0.979 | [0.967, 0.987] | 0.427      | +0.552      | N/A      |
| Qwen2.5-1.5B       | random_nouns            | L20   | 0.986 | [0.976, 0.992] | 0.426      | +0.560      | +0.007   |
| Qwen2.5-1.5B       | single_chars            | L11   | 0.957 | [0.936, 0.969] | 0.423      | +0.534      | -0.022   |
| Qwen2.5-1.5B       | all_same                | L0    | 1.000 | [1.000, 1.000] | 0.440      | +0.560      | +0.021   |
| Qwen2.5-1.5B       | numeric_vars            | L21   | 0.965 | [0.952, 0.972] | 0.422      | +0.543      | -0.014   |
| Qwen2.5-1.5B       | misleading_class_struct | L7    | 0.976 | [0.962, 0.982] | 0.416      | +0.559      | -0.004   |
| Qwen2.5-Coder-1.5B | baseline                | L8    | 0.982 | [0.969, 0.989] | 0.429      | +0.553      | N/A      |
| Qwen2.5-Coder-1.5B | random_nouns            | L19   | 0.978 | [0.965, 0.985] | 0.427      | +0.551      | -0.004   |
| Qwen2.5-Coder-1.5B | single_chars            | L7    | 0.956 | [0.936, 0.967] | 0.423      | +0.533      | -0.026   |
| Qwen2.5-Coder-1.5B | all_same                | L0    | 1.000 | [1.000, 1.000] | 0.440      | +0.560      | +0.018   |
| Qwen2.5-Coder-1.5B | numeric_vars            | L21   | 0.968 | [0.955, 0.975] | 0.424      | +0.544      | -0.014   |
| Qwen2.5-Coder-1.5B | misleading_class_struct | L7    | 0.978 | [0.967, 0.984] | 0.415      | +0.562      | -0.004   |
| StarCoder2-7B      | baseline                | L5    | 0.981 | [0.967, 0.990] | 0.478      | +0.503      | N/A      |
| StarCoder2-7B      | random_nouns            | L6    | 0.976 | [0.962, 0.983] | 0.466      | +0.510      | -0.006   |
| StarCoder2-7B      | single_chars            | L4    | 0.957 | [0.921, 0.971] | 0.475      | +0.482      | -0.025   |
| StarCoder2-7B      | all_same                | L0    | 1.000 | [1.000, 1.000] | 0.445      | +0.555      | +0.019   |
| StarCoder2-7B      | numeric_vars            | L23   | 0.968 | [0.956, 0.975] | 0.451      | +0.517      | -0.013   |
| StarCoder2-7B      | misleading_class_struct | L5    | 0.985 | [0.976, 0.990] | 0.467      | +0.518      | +0.004   |






Figure 1: Boolean probe and model-free baselines for Qwen2.5-1.5B across the four committed languages.

Table 8: Class/structure cross-language results. Transfer uses the Python-selected layer without a shared source-target problem partition. In-domain layers are selected separately. Only languages present in the committed exports are shown.

| model              | language   | programs | in-domain layer | in-domain F1 | transfer accuracy | transfer F1 |
|--------------------|------------|----------|-----------------|--------------|-------------------|-------------|
| Qwen2.5-1.5B       | Python     | 400      | L18             | 0.979        | N/A               | N/A         |
| Qwen2.5-1.5B       | C++        | 606      | L10             | 0.991        | 0.993             | 0.933       |
| Qwen2.5-1.5B       | JavaScript | 285      | L5              | 0.987        | 0.998             | 0.957       |
| Qwen2.5-1.5B       | C          | 97       | L8              | 0.986        | 0.984             | 0.886       |
| Qwen2.5-Coder-1.5B | Python     | 400      | L8              | 0.982        | N/A               | N/A         |
| Qwen2.5-Coder-1.5B | C++        | 606      | L8              | 0.988        | 0.991             | 0.918       |
| Qwen2.5-Coder-1.5B | JavaScript | 285      | L6              | 0.985        | 0.996             | 0.913       |
| Qwen2.5-Coder-1.5B | C          | 97       | L20             | 0.980        | 0.983             | 0.885       |
| StarCoder2-7B      | Python     | 400      | L5              | 0.981        | N/A               | N/A         |
| StarCoder2-7B      | C++        | 606      | L5              | 0.992        | 0.996             | 0.965       |
| StarCoder2-7B      | JavaScript | 285      | L3              | 0.989        | 0.999             | 0.980       |
| StarCoder2-7B      | C          | 97       | L5              | 0.986        | 0.987             | 0.906       |




Figure 2: Boolean probe and model-free baselines for Qwen2.5-Coder-1.5B across the four committed languages.

Table 9: Layerwise class/structure robustness summaries. The F1 range is taken over all committed layers. Maximum cosine similarity is measured against the baseline activation direction for each renamed condition.

| model              | strategy                | minimum F1 | maximum F1 (layer) | maximum cosine (layer) |
|--------------------|-------------------------|------------|--------------------|------------------------|
| Qwen2.5-1.5B       | all_same                | 0.999      | 1.000 (L0)         | 0.105 (L11)            |
| Qwen2.5-1.5B       | baseline                | 0.806      | 0.983 (L18)        | N/A                    |
| Qwen2.5-1.5B       | misleading_class_struct | 0.708      | 0.980 (L7)         | 0.311 (L15)            |
| Qwen2.5-1.5B       | numeric_vars            | 0.445      | 0.966 (L21)        | 0.329 (L22)            |
| Qwen2.5-1.5B       | random_nouns            | 0.529      | 0.989 (L15)        | 0.493 (L26)            |
| Qwen2.5-1.5B       | single_chars            | 0.530      | 0.965 (L7)         | 0.392 (L15)            |
| Qwen2.5-Coder-1.5B | all_same                | 0.999      | 1.000 (L0)         | 0.126 (L28)            |
| Qwen2.5-Coder-1.5B | baseline                | 0.807      | 0.985 (L7)         | N/A                    |
| Qwen2.5-Coder-1.5B | misleading_class_struct | 0.708      | 0.978 (L7)         | 0.324 (L7)             |
| Qwen2.5-Coder-1.5B | numeric_vars            | 0.445      | 0.969 (L24)        | 0.352 (L22)            |
| Qwen2.5-Coder-1.5B | random_nouns            | 0.530      | 0.983 (L23)        | 0.464 (L19)            |
| Qwen2.5-Coder-1.5B | single_chars            | 0.530      | 0.959 (L7)         | 0.415 (L15)            |
| StarCoder2-7B      | all_same                | 0.999      | 1.000 (L0)         | 0.101 (L21)            |
| StarCoder2-7B      | baseline                | 0.763      | 0.984 (L6)         | N/A                    |
| StarCoder2-7B      | misleading_class_struct | 0.717      | 0.985 (L3)         | 0.387 (L3)             |
| StarCoder2-7B      | numeric_vars            | 0.443      | 0.971 (L26)        | 0.369 (L22)            |
| StarCoder2-7B      | random_nouns            | 0.502      | 0.979 (L8)         | 0.511 (L3)             |
| StarCoder2-7B      | single_chars            | 0.529      | 0.962 (L3)         | 0.467 (L6)             |




Figure 3: Boolean probe and model-free baselines for StarCoder2-7B across the four committed languages.

Table 10: Complete iterator perturbation results under the shared five-seed protocol. Each condition refits a probe and selects its own layer on validation data. Role-conditioned and non-bijective strategies are confounded diagnostics rather than robustness tests.

| model              | strategy            | layer | F1    | 95% CI         | control F1 | selectivity | $\Delta$ |
|--------------------|---------------------|-------|-------|----------------|------------|-------------|----------|
| Qwen2.5-1.5B       | baseline            | L7    | 0.976 | [0.964, 0.982] | 0.424      | +0.552      | N/A      |
| Qwen2.5-1.5B       | random_nouns        | L11   | 0.961 | [0.951, 0.967] | 0.423      | +0.538      | -0.015   |
| Qwen2.5-1.5B       | single_chars        | L9    | 0.906 | [0.896, 0.917] | 0.428      | +0.479      | -0.069   |
| Qwen2.5-1.5B       | all_same            | L1    | 1.000 | [1.000, 1.000] | 0.456      | +0.544      | +0.024   |
| Qwen2.5-1.5B       | numeric_vars        | L23   | 0.910 | [0.899, 0.920] | 0.440      | +0.471      | -0.065   |
| Qwen2.5-1.5B       | misleading_iterator | L0    | 0.999 | [0.998, 0.999] | 0.403      | +0.595      | +0.023   |
| Qwen2.5-Coder-1.5B | baseline            | L7    | 0.978 | [0.968, 0.983] | 0.429      | +0.549      | N/A      |
| Qwen2.5-Coder-1.5B | random_nouns        | L12   | 0.956 | [0.946, 0.962] | 0.423      | +0.533      | -0.022   |
| Qwen2.5-Coder-1.5B | single_chars        | L8    | 0.906 | [0.895, 0.916] | 0.426      | +0.480      | -0.072   |
| Qwen2.5-Coder-1.5B | all_same            | L1    | 1.000 | [1.000, 1.000] | 0.459      | +0.541      | +0.022   |
| Qwen2.5-Coder-1.5B | numeric_vars        | L22   | 0.911 | [0.896, 0.921] | 0.440      | +0.470      | -0.068   |
| Qwen2.5-Coder-1.5B | misleading_iterator | L0    | 0.998 | [0.997, 0.999] | 0.402      | +0.596      | +0.020   |
| StarCoder2-7B      | baseline            | L11   | 0.987 | [0.981, 0.990] | 0.461      | +0.526      | N/A      |
| StarCoder2-7B      | random_nouns        | L6    | 0.975 | [0.967, 0.979] | 0.455      | +0.520      | -0.012   |
| StarCoder2-7B      | single_chars        | L6    | 0.941 | [0.933, 0.949] | 0.458      | +0.483      | -0.046   |
| StarCoder2-7B      | all_same            | L2    | 1.000 | [1.000, 1.000] | 0.460      | +0.540      | +0.013   |
| StarCoder2-7B      | numeric_vars        | L26   | 0.934 | [0.921, 0.941] | 0.454      | +0.480      | -0.053   |
| StarCoder2-7B      | misleading_iterator | L0    | 0.998 | [0.997, 0.999] | 0.406      | +0.592      | +0.012   |




Figure 4: Boolean probe layer curves per language across the three models.




Figure 5: Problem-paired boolean refit-performance deltas on Python. No corresponding committed renaming run exists for other languages.

Table 11: Model-free iterator surface baselines on the same 427 Python programs. Scores are occurrence-level macro-F1 and are not a paired probe-minus-surface interval.

| comparator       | macro-F1 |
|------------------|----------|
| majority         | 0.363    |
| covariates       | 0.361    |
| masked line      | 0.844    |
| masked statement | 0.842    |
| masked window    | 0.869    |
| name only        | 0.918    |

Table 12: Complete iterator cross-language transfer results. Transfer columns use the Python-selected layer.

| model              | language   | programs | in-domain layer | in-domain F1 | transfer accuracy | transfer F1 |
|--------------------|------------|----------|-----------------|--------------|-------------------|-------------|
| Qwen2.5-1.5B       | Python     | 256      | L9              | 0.971        | N/A               | N/A         |
| Qwen2.5-1.5B       | Java       | 188      | L13             | 0.990        | 0.993             | 0.956       |
| Qwen2.5-1.5B       | C++        | 191      | L12             | 0.992        | 0.994             | 0.965       |
| Qwen2.5-1.5B       | C          | 191      | L7              | 0.972        | 0.987             | 0.934       |
| Qwen2.5-1.5B       | C#         | 189      | L13             | 0.987        | 0.993             | 0.959       |
| Qwen2.5-1.5B       | JavaScript | 266      | L9              | 0.985        | 0.992             | 0.966       |
| Qwen2.5-1.5B       | PHP        | 253      | L12             | 0.986        | 0.996             | 0.981       |
| Qwen2.5-Coder-1.5B | Python     | 256      | L10             | 0.978        | N/A               | N/A         |
| Qwen2.5-Coder-1.5B | Java       | 188      | L11             | 0.989        | 0.986             | 0.908       |
| Qwen2.5-Coder-1.5B | C++        | 191      | L12             | 0.988        | 0.988             | 0.923       |
| Qwen2.5-Coder-1.5B | C          | 191      | L6              | 0.967        | 0.982             | 0.898       |
| Qwen2.5-Coder-1.5B | C#         | 189      | L14             | 0.984        | 0.985             | 0.909       |
| Qwen2.5-Coder-1.5B | JavaScript | 266      | L10             | 0.982        | 0.983             | 0.915       |
| Qwen2.5-Coder-1.5B | PHP        | 253      | L12             | 0.988        | 0.986             | 0.923       |
| StarCoder2-7B      | Python     | 256      | L9              | 0.986        | N/A               | N/A         |
| StarCoder2-7B      | Java       | 188      | L7              | 0.992        | 0.980             | 0.887       |
| StarCoder2-7B      | C++        | 191      | L6              | 0.989        | 0.984             | 0.906       |
| StarCoder2-7B      | C          | 191      | L10             | 0.971        | 0.981             | 0.901       |
| StarCoder2-7B      | C#         | 189      | L8              | 0.987        | 0.980             | 0.889       |
| StarCoder2-7B      | JavaScript | 266      | L7              | 0.989        | 0.980             | 0.911       |
| StarCoder2-7B      | PHP        | 253      | L12             | 0.989        | 0.987             | 0.938       |

Table 13: Layerwise iterator robustness summaries. F1 ranges cover every committed layer. Cosine similarity compares each perturbation direction with the baseline direction.

| model              | strategy            | minimum F1 | maximum F1 (layer) | maximum cosine (layer) |
|--------------------|---------------------|------------|--------------------|------------------------|
| Qwen2.5-1.5B       | all_same            | 0.999      | 1.000 (L0)         | 0.200 (L15)            |
| Qwen2.5-1.5B       | baseline            | 0.893      | 0.979 (L11)        | N/A                    |
| Qwen2.5-1.5B       | misleading_iterator | 0.988      | 0.999 (L8)         | 0.129 (L7)             |
| Qwen2.5-1.5B       | numeric_vars        | 0.550      | 0.913 (L23)        | 0.255 (L19)            |
| Qwen2.5-1.5B       | random_nouns        | 0.628      | 0.964 (L12)        | 0.392 (L19)            |
| Qwen2.5-1.5B       | single_chars        | 0.658      | 0.918 (L7)         | 0.265 (L21)            |
| Qwen2.5-Coder-1.5B | all_same            | 0.999      | 1.000 (L0)         | 0.227 (L18)            |
| Qwen2.5-Coder-1.5B | baseline            | 0.893      | 0.979 (L10)        | N/A                    |
| Qwen2.5-Coder-1.5B | misleading_iterator | 0.979      | 0.999 (L4)         | 0.179 (L20)            |
| Qwen2.5-Coder-1.5B | numeric_vars        | 0.550      | 0.913 (L23)        | 0.277 (L19)            |
| Qwen2.5-Coder-1.5B | random_nouns        | 0.628      | 0.961 (L16)        | 0.448 (L14)            |
| Qwen2.5-Coder-1.5B | single_chars        | 0.658      | 0.912 (L8)         | 0.299 (L17)            |
| StarCoder2-7B      | all_same            | 0.999      | 1.000 (L0)         | 0.169 (L20)            |
| StarCoder2-7B      | baseline            | 0.890      | 0.988 (L6)         | N/A                    |
| StarCoder2-7B      | misleading_iterator | 0.990      | 0.998 (L0)         | 0.234 (L21)            |
| StarCoder2-7B      | numeric_vars        | 0.549      | 0.937 (L24)        | 0.271 (L25)            |
| StarCoder2-7B      | random_nouns        | 0.603      | 0.975 (L5)         | 0.385 (L10)            |
| StarCoder2-7B      | single_chars        | 0.657      | 0.946 (L5)         | 0.359 (L12)            |




Figure 6: Iterator probe macro-F1 across layers on original names, with the name-only surface baseline marked.

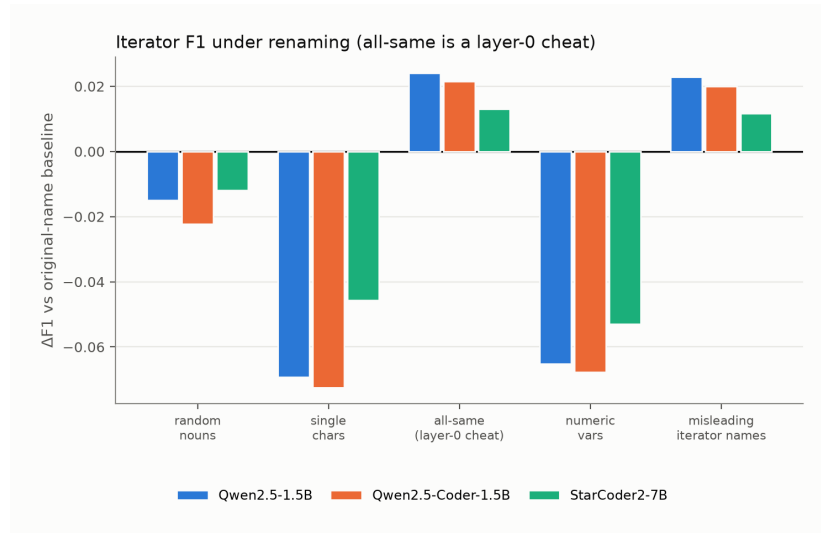


Figure 7: Iterator perturbation deltas relative to baseline.




Figure 8: Iterator cosine similarity to the baseline direction.




Figure 9: Available iterator cross-language transfer results.




Figure 10: Iterator probe versus matched model-free surface baselines on the same 427 Python programs.




Figure 11: Iterator probe macro-F1 across layers and renaming strategies.




Figure 12: Iterator val-selected F1 by renaming strategy, with BCa 95% intervals.




Figure 13: Iterator probe F1 by strategy and layer.






Figure 14: Exploratory iterator activation-patching recovery under Python name perturbations. The terminal layer is one by construction.




Figure 15: Exploratory iterator activation-patching recovery under cross-language transfer. C-language rows with zero pairs are omitted.

## 287 B Causal intervention details

Table 14: Complete gate-specified class/structure patching summary in Qwen2.5-1.5B. The causal gate required mean effect at least 0.10, a positive interval, separation from controls, and recovery at least 0.05 in both directions.

| layer | span       | direction | control     | <i>n</i> | mean effect | 95% CI             | recovery | minus random |
|-------|------------|-----------|-------------|----------|-------------|--------------------|----------|--------------|
| 0     | query_name | denoise   | target      | 288      | +0.0000     | [0.0000, 0.0000]   | +0.0000  | N/A          |
| 0     | query_name | noise     | target      | 288      | +0.0000     | [0.0000, 0.0000]   | +0.0000  | N/A          |
| 18    | placebo    | denoise   | target      | 288      | +0.0015     | [-0.0010, 0.0043]  | +0.0015  | N/A          |
| 18    | placebo    | noise     | target      | 288      | +0.0006     | [-0.0019, 0.0032]  | +0.0006  | N/A          |
| 18    | query_name | denoise   | random      | 288      | -0.0107     | [-0.0205, -0.0014] | -0.0110  | N/A          |
| 18    | query_name | denoise   | same_source | 288      | +0.0004     | [-0.0009, 0.0016]  | +0.0004  | N/A          |
| 18    | query_name | denoise   | target      | 288      | +0.0087     | [-0.0007, 0.0188]  | +0.0089  | +0.0194      |
| 18    | query_name | noise     | random      | 288      | +0.0096     | [0.0023, 0.0164]   | +0.0099  | N/A          |
| 18    | query_name | noise     | same_source | 288      | -0.0005     | [-0.0018, 0.0006]  | -0.0006  | N/A          |
| 18    | query_name | noise     | target      | 288      | +0.0199     | [0.0109, 0.0296]   | +0.0204  | +0.0103      |

288 **Status.** The iterator recovery exports predate the gate-specified class/structure protocol. They  
 289 report descriptive recovery curves without clustered confidence intervals or matched placebo/random  
 290 controls. The terminal value is one by construction, and C-language cross-language rows with zero  
 291 pairs are missing data rather than null effects.

Table 15: Aggregate probe-link diagnostic for the class/structure evaluation pairs. Absolute values are reported because pair orientation is arbitrary. Large probe movement alongside small behavioral movement is descriptive of the failed causal transfer, not evidence that probe movement causes behavior.

| quantity            | mean absolute | median absolute |
|---------------------|---------------|-----------------|
| probe movement      | 15.786        | 16.302          |
| behavioral movement | 0.038         | 0.031           |

Table 16: Class/structure gate and readout audit. Both prompt classes favor True. The pairwise gap remains separable, but the binary readout is not calibrated to distinguish function prompts.

| diagnostic                  | estimate    | 95% CI         | gate |
|-----------------------------|-------------|----------------|------|
| class logit difference      | 2.752       | [2.624, 2.898] | pass |
| function logit difference   | 1.778       | [1.652, 1.909] | fail |
| class-minus-function gap    | 0.974       | [0.908, 1.035] | pass |
| pair-gap accuracy           | 0.986       | [0.962, 0.993] | pass |
| query/declaration probe AUC | 0.943/0.905 | N/A            | pass |

Table 17: Patching audit trail. Completeness establishes that the scheduled Qwen evaluation finished. It does not turn the failed causal gate into positive evidence. The controller stopped before Coder and StarCoder2 after the Qwen null.

| audit item                           | value                |
|--------------------------------------|----------------------|
| behavior rows present/expected       | 1152/1152            |
| primary rows present/expected        | 4032/4032            |
| baseline probe-gap/behavior Spearman | 0.164 [0.023, 0.286] |
| controller terminal state            | stopped_qwen_null    |

Table 18: Exploratory iterator activation-patching recovery for Qwen2.5-1.5B, within-language perturbations.

| condition           | readout | pairs | clean/corrupt | L1    | L2    | L3    | L4    | L5    | L6    | L7    |
|---------------------|---------|-------|---------------|-------|-------|-------|-------|-------|-------|-------|
| random_nouns        | L7      | 150   | 0.999/0.307   | 0.993 | 0.963 | 0.994 | 0.994 | 0.995 | 0.999 | 1.000 |
| single_chars        | L7      | 150   | 0.995/0.712   | 0.580 | 0.530 | 0.820 | 0.927 | 0.964 | 0.994 | 1.000 |
| all_same            | L7      | 150   | 0.999/0.138   | 0.321 | 0.454 | 0.879 | 0.955 | 0.984 | 1.000 | 1.000 |
| numeric_vars        | L7      | 150   | 0.999/0.211   | 0.949 | 0.927 | 0.976 | 0.986 | 0.989 | 1.000 | 1.000 |
| misleading_iterator | L7      | 150   | 0.998/0.324   | 0.869 | 0.900 | 0.970 | 0.985 | 0.988 | 1.000 | 1.000 |

Table 19: Exploratory iterator activation-patching recovery for Qwen2.5-1.5B, cross-language transfer.

| condition  | readout | pairs | clean/corrupt | L1     | L2     | L3    | L4    | L5    | L6    | L7    | L8    | L9    |
|------------|---------|-------|---------------|--------|--------|-------|-------|-------|-------|-------|-------|-------|
| Java       | L9      | 20    | 0.998/0.768   | -0.409 | -0.044 | 0.165 | 0.416 | 0.767 | 0.625 | 0.544 | 0.505 | 1.000 |
| C++        | L9      | 20    | 0.996/0.801   | -0.790 | -0.345 | 0.010 | 0.148 | 0.544 | 0.545 | 0.486 | 0.275 | 1.000 |
| C          | L9      | 0     | 0.000/0.000   | N/A    | N/A    | N/A   | N/A   | N/A   | N/A   | N/A   | N/A   | N/A   |
| C#         | L9      | 30    | 0.993/0.801   | 0.059  | 0.374  | 0.514 | 0.657 | 0.809 | 0.734 | 0.680 | 0.681 | 1.000 |
| JavaScript | L9      | 44    | 0.996/0.828   | -0.011 | 0.128  | 0.465 | 0.734 | 0.827 | 0.773 | 0.745 | 0.774 | 1.000 |
| PHP        | L9      | 12    | 0.987/0.810   | 0.427  | 0.437  | 0.521 | 0.571 | 0.586 | 0.325 | 0.353 | 0.437 | 1.000 |

Table 20: Exploratory iterator activation-patching recovery for Qwen2.5-Coder-1.5B, within-language perturbations.

| condition           | readout | pairs | clean/corrupt | L1    | L2    | L3    | L4    | L5    | L6    |
|---------------------|---------|-------|---------------|-------|-------|-------|-------|-------|-------|
| random_nouns        | L6      | 150   | 0.999/0.138   | 0.991 | 0.930 | 0.994 | 0.996 | 0.998 | 1.000 |
| single_chars        | L6      | 150   | 0.996/0.614   | 0.860 | 0.799 | 0.935 | 0.963 | 0.983 | 1.000 |
| all_same            | L6      | 150   | 0.999/0.142   | 0.808 | 0.900 | 0.961 | 0.986 | 0.998 | 1.000 |
| numeric_vars        | L6      | 150   | 0.999/0.057   | 0.966 | 0.961 | 0.991 | 0.997 | 0.999 | 1.000 |
| misleading_iterator | L6      | 150   | 0.999/0.286   | 0.820 | 0.835 | 0.916 | 0.946 | 0.983 | 1.000 |

Table 21: Exploratory iterator activation-patching recovery for Qwen2.5-Coder-1.5B, cross-language transfer.

| condition  | readout | pairs | clean/corrupt | L1     | L2     | L3    | L4    | L5    | L6    | L7    |
|------------|---------|-------|---------------|--------|--------|-------|-------|-------|-------|-------|
| Java       | L7      | 10    | 0.989/0.709   | -0.300 | -0.146 | 0.616 | 0.867 | 0.845 | 0.940 | 1.000 |
| C++        | L7      | 11    | 0.979/0.769   | -0.066 | 0.025  | 0.785 | 0.869 | 0.869 | 0.904 | 1.000 |
| C          | L7      | 0     | 0.000/0.000   | N/A    | N/A    | N/A   | N/A   | N/A   | N/A   | N/A   |
| C#         | L7      | 9     | 0.981/0.629   | 0.307  | 0.353  | 0.772 | 0.867 | 0.903 | 0.989 | 1.000 |
| JavaScript | L7      | 13    | 0.983/0.766   | 0.708  | 0.678  | 0.968 | 0.943 | 0.953 | 0.983 | 1.000 |
| PHP        | L7      | 7     | 0.982/0.748   | -0.211 | -0.397 | 0.614 | 0.760 | 0.934 | 1.027 | 1.000 |

Table 22: Exploratory iterator activation-patching recovery for StarCoder2-7B, within-language perturbations.

| condition           | readout | pairs | clean/corrupt | L1    | L2    | L3    | L4    | L5    |
|---------------------|---------|-------|---------------|-------|-------|-------|-------|-------|
| random_nouns        | L5      | 150   | 0.998/0.569   | 0.931 | 0.766 | 0.966 | 0.951 | 1.000 |
| single_chars        | L5      | 150   | 0.994/0.686   | 0.636 | 0.006 | 0.823 | 0.759 | 1.000 |
| all_same            | L5      | 150   | 0.999/0.249   | 0.342 | 0.772 | 0.946 | 0.979 | 1.000 |
| numeric_vars        | L5      | 150   | 0.998/0.530   | 0.889 | 0.811 | 0.953 | 0.955 | 1.000 |
| misleading_iterator | L5      | 150   | 0.997/0.628   | 0.777 | 0.104 | 0.921 | 0.843 | 1.000 |

Table 23: Exploratory iterator activation-patching recovery for StarCoder2-7B, cross-language transfer.

| condition  | readout | pairs | clean/corrupt | L1    | L2    | L3    | L4    | L5    | L6    | L7    | L8    | L9    | L10   |
|------------|---------|-------|---------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| Java       | L10     | 150   | 0.996/0.750   | 0.254 | 0.634 | 0.829 | 0.913 | 0.957 | 0.964 | 0.979 | 0.988 | 0.994 | 1.000 |
| C++        | L10     | 150   | 0.996/0.771   | 0.337 | 0.756 | 0.892 | 0.938 | 0.966 | 0.969 | 0.979 | 0.988 | 0.995 | 1.000 |
| C          | L10     | 0     | 0.000/0.000   | N/A   | N/A   | N/A   | N/A   | N/A   | N/A   | N/A   | N/A   | N/A   | N/A   |
| C#         | L10     | 150   | 0.996/0.740   | 0.294 | 0.692 | 0.877 | 0.942 | 0.969 | 0.974 | 0.984 | 0.991 | 0.996 | 1.000 |
| JavaScript | L10     | 150   | 0.996/0.716   | 0.311 | 0.743 | 0.880 | 0.939 | 0.972 | 0.976 | 0.983 | 0.990 | 0.996 | 1.000 |
| PHP        | L10     | 73    | 0.993/0.808   | 0.491 | 0.796 | 0.938 | 0.964 | 0.979 | 0.973 | 0.983 | 0.994 | 0.999 | 1.000 |

## C Exploratory boolean interventions

**Status and estimand.** These boolean interventions are exploratory diagnostics rather than evidence for the main causal claim. The committed export contains layer-wise means but not the per-case arrays needed to recompute clustered intervals, and the runs do not use the specified gate of the class/structure experiment. The exporter defines effect as  $|\text{clean} - \text{intervened}|$  and specificity as that effect divided by the corresponding random-position-control effect. Specificity can become large when its denominator is small. We therefore report both specificity at the largest-effect layer and the maximum specificity over layers rather than selecting only the more favorable one.

Table 24: Exploratory boolean patch diagnostics. Effect and specificity use the exporter definitions above. No uncertainty interval is available from the committed summary.

| language   | model              | $n$ | peak layer | peak effect | spec. at peak | best spec. (layer) |
|------------|--------------------|-----|------------|-------------|---------------|--------------------|
| C++        | Qwen2.5-1.5B       | 157 | L0         | 14.845      | $3.6\times$   | $7.9\times$ (L20)  |
| C++        | Qwen2.5-Coder-1.5B | 158 | L0         | 15.414      | $3.6\times$   | $6.9\times$ (L8)   |
| C++        | StarCoder2-7B      | 157 | L4         | 14.982      | $8.1\times$   | $10.4\times$ (L16) |
| Java       | Qwen2.5-1.5B       | 155 | L8         | 11.607      | $6.1\times$   | $7.7\times$ (L12)  |
| Java       | Qwen2.5-Coder-1.5B | 155 | L8         | 12.642      | $5.9\times$   | $7.7\times$ (L12)  |
| Java       | StarCoder2-7B      | 155 | L4         | 13.900      | $7.5\times$   | $8.6\times$ (L12)  |
| JavaScript | Qwen2.5-1.5B       | 58  | L8         | 12.755      | $6.9\times$   | $6.9\times$ (L8)   |
| JavaScript | Qwen2.5-Coder-1.5B | 62  | L8         | 12.717      | $6.6\times$   | $7.4\times$ (L12)  |
| JavaScript | StarCoder2-7B      | 129 | L4         | 12.137      | $8.1\times$   | $8.2\times$ (L16)  |
| PHP        | Qwen2.5-1.5B       | 15  | L4         | 11.698      | $7.2\times$   | $7.2\times$ (L4)   |
| PHP        | Qwen2.5-Coder-1.5B | 16  | L4         | 11.425      | $4.7\times$   | $7.4\times$ (L0)   |
| PHP        | StarCoder2-7B      | 17  | L4         | 11.545      | $6.3\times$   | $9.3\times$ (L16)  |
| Python     | Qwen2.5-1.5B       | 176 | L8         | 12.948      | $6.0\times$   | $7.2\times$ (L16)  |
| Python     | Qwen2.5-Coder-1.5B | 176 | L8         | 13.275      | $6.7\times$   | $7.8\times$ (L16)  |
| Python     | StarCoder2-7B      | 176 | L4         | 13.592      | $7.9\times$   | $9.7\times$ (L12)  |

Table 25: Exploratory boolean steer diagnostics. Effect and specificity use the exporter definitions above. No uncertainty interval is available from the committed summary.

| language   | model              | $n$ | peak layer | peak effect | spec. at peak  | best spec. (layer)   |
|------------|--------------------|-----|------------|-------------|----------------|----------------------|
| C++        | Qwen2.5-1.5B       | 79  | L24        | 2.508       | $45.0\times$   | $169.3\times$ (L4)   |
| C++        | Qwen2.5-Coder-1.5B | 79  | L24        | 1.436       | $602.0\times$  | $602.0\times$ (L24)  |
| C++        | StarCoder2-7B      | 78  | L20        | 0.653       | $6.9\times$    | $345.4\times$ (L4)   |
| Java       | Qwen2.5-1.5B       | 78  | L24        | 2.486       | $72.2\times$   | $826.9\times$ (L8)   |
| Java       | Qwen2.5-Coder-1.5B | 78  | L24        | 1.712       | $72.0\times$   | $179.9\times$ (L0)   |
| Java       | StarCoder2-7B      | 78  | L24        | 0.555       | $244.0\times$  | $244.0\times$ (L24)  |
| JavaScript | Qwen2.5-1.5B       | 31  | L24        | 1.763       | $24.5\times$   | $76.8\times$ (L8)    |
| JavaScript | Qwen2.5-Coder-1.5B | 31  | L24        | 1.147       | $65.3\times$   | $147.6\times$ (L4)   |
| JavaScript | StarCoder2-7B      | 65  | L16        | 0.591       | $11.2\times$   | $61.0\times$ (L28)   |
| PHP        | Qwen2.5-1.5B       | 7   | L24        | 2.759       | $190.2\times$  | $190.2\times$ (L24)  |
| PHP        | Qwen2.5-Coder-1.5B | 8   | L24        | 0.898       | $11.5\times$   | $132.7\times$ (L8)   |
| PHP        | StarCoder2-7B      | 9   | L20        | 0.991       | $8.7\times$    | $99.9\times$ (L4)    |
| Python     | Qwen2.5-1.5B       | 88  | L24        | 2.175       | $209.1\times$  | $794.6\times$ (L16)  |
| Python     | Qwen2.5-Coder-1.5B | 88  | L24        | 1.231       | $1027.1\times$ | $1027.1\times$ (L24) |
| Python     | StarCoder2-7B      | 88  | L20        | 0.721       | $8.6\times$    | $55.1\times$ (L4)    |

**Omitted figures.** Per-cell diagnostic plots for the exploratory boolean interventions and the committed role analyses are reproducible from the same artifacts with `make_interp_appendix.py -full-figures`. They visualize the tables above and add no uncertainty estimate the exports lack.

Table 26: Exploratory boolean ablate diagnostics. Effect and specificity use the exporter definitions above. No uncertainty interval is available from the committed summary.

| language   | model              | $n$ | peak layer | peak effect | spec. at peak | best spec. (layer) |
|------------|--------------------|-----|------------|-------------|---------------|--------------------|
| C++        | Qwen2.5-1.5B       | 157 | L0         | 14.286      | 4.5×          | 31.9× (L4)         |
| C++        | Qwen2.5-Coder-1.5B | 158 | L0         | 12.927      | 4.4×          | 76.2× (L20)        |
| C++        | StarCoder2-7B      | 157 | L0         | 13.405      | 4.3×          | 36.5× (L4)         |
| Java       | Qwen2.5-1.5B       | 155 | L16        | 9.720       | 28.5×         | 86.0× (L12)        |
| Java       | Qwen2.5-Coder-1.5B | 155 | L4         | 9.331       | 12.3×         | 62.4× (L12)        |
| Java       | StarCoder2-7B      | 155 | L4         | 11.190      | 27.0×         | 27.0× (L4)         |
| JavaScript | Qwen2.5-1.5B       | 58  | L16        | 10.852      | 26.7×         | 26.7× (L16)        |
| JavaScript | Qwen2.5-Coder-1.5B | 62  | L16        | 9.868       | 28.1×         | 116.4× (L24)       |
| JavaScript | StarCoder2-7B      | 129 | L4         | 10.970      | 27.3×         | 27.3× (L4)         |
| PHP        | Qwen2.5-1.5B       | 15  | L16        | 15.698      | 15.0×         | 102.4× (L20)       |
| PHP        | Qwen2.5-Coder-1.5B | 16  | L16        | 13.402      | 26.5×         | 79.2× (L20)        |
| PHP        | StarCoder2-7B      | 17  | L4         | 14.260      | 19.4×         | 19.7× (L8)         |
| Python     | Qwen2.5-1.5B       | 176 | L0         | 12.144      | 5.8×          | 36.7× (L24)        |
| Python     | Qwen2.5-Coder-1.5B | 176 | L0         | 11.286      | 6.8×          | 33.1× (L16)        |
| Python     | StarCoder2-7B      | 176 | L4         | 11.818      | 26.5×         | 26.5× (L4)         |

## D Explicit coverage holes

- Accumulator and index/key results use legacy exports. No committed problem-grouped, five-seed CSV reproduces those roles under the modern protocol.
- Iterator has a model-free surface comparator. The probe-minus-surface gap is not a paired clustered interval.
- Class/structure transfer exports include Python, C++, JavaScript, and C, but not Java, C#, or PHP.
- Boolean probe exports include Python, Java, JavaScript, and PHP. C, C#, and C++ are absent, and renaming is available only for Python.
- Gate-specified class/structure patching stopped after the Qwen2.5-1.5B null. Coder, StarCoder2, and the all-layer core sweep were not attempted.
- No matched model-free surface comparator is committed for the class/structure role.

## E Reproducibility

The appendix is generated from committed aggregate boolean, iterator, class or structure, and patching artifacts with `uv run python scripts/make_interp_appendix.py`. The claim audit uses `uv run python tests/test_interp4d_claims.py`. Modern probes use problem hashes, validation-selected layers, and five seeds where stated. The class or structure intervention records model and dataset revisions, hashes, all 4,032 rows across its primary and behavior schedules, and clustered intervals. The paired boolean CSV is aggregate evidence because its aligned prediction inputs were not retained. The regression test checks manuscript-to-artifact consistency rather than reproducing activation extraction or probe fitting.

## F Broader impacts

Positive effects are scientific. Probe scores near 0.98 are easy to read as evidence that a model computed a role. If such claims later inform audits of coding assistants, an outcome-aware record reduces that overclaim.

Negative effects follow from the same scores. A reporting matrix can be treated as permission to stop at decodability if the bound is ignored. Role labels derived from public competitive-programming solutions could be reused to train name-based heuristics that recover the label from the identifier. No new generative model is released. The source programs were already public.

## G Existing assets and licenses

Source programs come from XLCOST under Apache 2.0 [Zhu et al., 2022]. Qwen2.5-1.5B and Qwen2.5-Coder-1.5B are used under Apache 2.0. StarCoder2-7B is used under the BigCode OpenRAIL-M license. Each release is cited by name and identifier. Derived role labels inherit Apache 2.0 from XLCOST.

## H New assets

The new assets are structurally labeled XLCOST variable-role data and the analysis scripts that produce the reported tables. A dataset card documents fields, extractors, splits, and the parent license. Committed result files carry artifact provenance. A named download URL is omitted from this double-blind PDF.

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction distinguish decodability, surface sufficiency, intervention confounding, and site-state causal relevance. They describe heterogeneous identifier properties and roles across three models and seven languages collectively without claiming a complete factorial design.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The Limitations section covers uneven coverage, heterogeneous targets, missing extractor audits, role-conditioned and non-bijective renaming, cross-language overlap, comparator selection, missing paired inputs, and the one-model, one-site causal gate.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.

- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper reports no theoretical results.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The study-design section specifies prediction units, controls, split units, layer selection, and seeds. The reproducibility appendix records how tables are generated from committed artifacts, including which boolean comparison is aggregate because aligned prediction inputs were not retained.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed

instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.

- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Analysis scripts, committed result files, and the labeled dataset are available with the project. The reproducibility appendix names the generator and the claim-audit command. A named repository URL is omitted from this double-blind PDF.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The study-design section specifies prediction units, controls, split units, layer selection, and seeds. The reproducibility appendix records how tables are generated from committed artifacts.



Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The primary boolean comparison and class or structure intervention report problem-clustered or template-clustered 95% intervals. Iterator reports BCa intervals on validation-selected F1. Legacy accumulator and index observations lack comparable intervals and are labeled exploratory rather than confirmatory.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: The paper does not report hardware, memory, or complete wall-clock costs.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work analyses publicly released models on a public benchmark of competitive-programming solutions. It involves no human subjects, no personally identifying data, and no deployment.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Appendix F states a scientific benefit for more conservative probe claims and a misuse path if the matrix is treated as permission to stop at decodability or if role labels train name-based heuristics.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The labeled dataset is a derivative of a public competitive-programming benchmark. It does not release a high-risk generative model or scraped media that would require extra safeguards.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Appendix G names Apache 2.0 for XLCoST, Qwen2.5-1.5B, and Qwen2.5-Coder-1.5B, and the BigCode OpenRAIL-M license for StarCoder2-7B. Derived labels inherit Apache 2.0. Each asset is cited.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

## 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: Appendix H documents the labeled XLCoST variable-role data, the dataset card covering fields and the parent license, and committed analysis artifacts. A named URL is omitted from this double-blind PDF.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

## 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper involves no crowdsourcing and no human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

#### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper involves no human subjects, so no IRB review was required.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

#### 16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: Large language models are the object of study, not an important or non-standard component used to develop the research method. Writing and editing assistance is exempt under the question.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.